

# PalmReader.ai: A Deep Network Driven Fortune-Teller

CHUNG, Pok Yu LAM, Ching Yin LI, Ho Ming

The Hong Kong University of Science and Technology  
Clear Water Bay, Hong Kong

pychungae@connect.ust.hk cylambu@connect.ust.hk hmliaf@connect.ust.hk

## Abstract

*Palm-reading is the practice of fortune-telling by reading one’s palm. While it dates back to centuries ago, it is regarded as a pseudoscience by many due to a lack of scientific evidence, and is generally performed subjectively by humans. However, with modern statistical and deep learning techniques, we propose an automatic palm reading system by learning characteristics and hidden features within datasets. We construct a pipeline with vision transformers, object detection, and large language models to produce palm reading results and predictions from simply an image of the user’s palm. Our project also uncovers the statistical connection between palm features and personal life experiences, creating a path for further related work.*

## 1. Introduction

Fortune-telling through palm reading, known as palmistry, has been practiced for centuries. However, such methods lack scientific support and are regarded by many as pseudoscience. There is also a lack of evidence for the connection between palm line features and the predictions stemming from this form of fortune-telling.

In the age of machine learning, large datasets can be leveraged along with machine learning techniques, such as deep neural networks, to uncover hidden characteristics and connections that were previously difficult to observe. Under this predicate, we propose that there may be statistical connections between palm features and palmistry fortune-telling results.

We construct a full pipeline to predict scores for aspects of the user’s life, such as character, emotional predispositions, health status, and life path. Our pipeline is composed of the following components.

- **Vision Transformer (ViT):** A vision transformer serves to extract features to vector embeddings that encode geometric features of palm lines in the input image.
- **Text decoder:** We use a text decoder to generate factual

descriptions of the user’s palm lines in natural language from the embeddings produced by the vision transformer.

- **Scoring system:** We score different aspects of life by examining the geometric features of the detected palm lines, using a neural network.
- **LLM inference:** We produce, from the description of palm lines and scores, a report in natural language that outputs fortune-telling content.

## 2. Problem Formulation

### 2.1. Four Primary Lines for Palm Line Analysis

Palm line analysis is an ancient practice that interprets the lines on the palm of the hand to glean insights into an individual’s personality traits, emotional tendencies, and life experiences. The four primary lines—Heart, Head, Life, and Fate—are central to this practice, each offering unique perspectives into various aspects of human life. All of our following tasks resolving these four lines [3]:

- **Heart Line:** Positioned at the top of the palm, the Heart Line is believed to reflect emotional stability, romantic perspectives, and cardiac health.
- **Head Line:** Running horizontally across the palm, the Head Line represents intellectualism, learning style, and communication abilities.
- **Life Line:** Encircling the thumb, the Life Line is traditionally associated with physical health, general well-being, and major life changes.
- **Fate Line:** Often running vertically from the base of the palm, the Fate Line indicates the degree to which external circumstances influence an individual’s life path.

Therefore, one of our primary objectives is to successfully segment these primary palm lines. Achieving precise segmentation is essential before we can undertake further analytical tasks. This foundational step ensures that subsequent analysis will be based on accurate and distinct representations of each line.

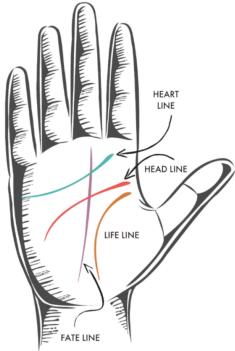


Figure 1. Different palm lines.

## 2.2. Feature Extraction and Text Decoder

Following the segmentation of the palm lines, a critical subsequent step involves utilizing a text decoder to meticulously analyze the shapes and other significant characteristics that may influence the analysis. This detailed examination is essential before the data is passed to a large language model for comprehensive interpretation and final analysis.

## 2.3. Quantified Reasoning

To enhance the reasoning process and mitigate any uncertainties, it is essential to implement a scoring system for the features identified in the palm lines. This scoring will assess various attributes, such as indicators of luck, health, and other relevant characteristics. By quantifying these features, we can provide a more structured and objective analysis that informs subsequent interpretations and enhances the overall reliability of the insights derived from the palm analysis.

## 2.4. Final Analysis

After obtaining scores for various attributes and elaborating on the text descriptions of the palm lines, the final and critical step is to generate a comprehensive, user-friendly analysis. This task will be accomplished using a large language model (LLM), which will synthesize the scores and textual insights into a coherent narrative. The LLM will provide an accessible interpretation of the analysis, enabling users to easily understand the implications of the palm features concerning aspects such as luck, health, and personal insights. By leveraging the capabilities of the LLM, we aim to deliver a polished and informative report that enhances the user experience and facilitates deeper engagement with the results of the palm analysis.

## 3. Dataset Design

### 3.1. Data Origin and Characteristics

The data utilized in this study is purely based on images of hands, sourced from the Kaggle dataset titled "Human Palm Images." [1] To ensure a diverse and representative dataset, we selected images that capture an overhead view of hands, which aligns with our objective of scanning users' hands effectively. The dataset is well-balanced, encompassing both left and right hands, as well as involving both male and female images of 400 each. This careful selection aims to mitigate potential biases and ethical concerns associated with dataset representation. By prioritizing diversity in our dataset, we strive to prevent discriminatory practices and ensure that our findings are applicable across different demographics. This approach not only enhances the robustness of our analysis but also upholds ethical standards in research.

### 3.2. Data Preprocessing

The data preprocessing pipeline for our AI application involves several critical steps to enhance the quality of the input images and prepare them for effective analysis. We perform the following operations on the dataset.

#### 3.2.1. Background Extraction

We isolate the palm from the surrounding background in an image. By effectively removing the background, we ensure that the model focuses exclusively on the palm, which leads to more accurate segmentation. We assume that the images are simple and there is no occlusion of our target objects.

#### 3.2.2. Transformation to Grayscale

Grayscale conversion is the process of transforming a color image into a single-channel format that represents varying shades of gray. This technique simplifies the complexity of the image by reducing it from three color channels (RGB) to one, making it easier to process [2]. In the context of palm segmentation, grayscale images enhance the visibility of edges and contours while eliminating color distractions that may not be relevant to the task. By focusing solely on intensity values, we can better highlight the structural features of the palm, which is essential for accurately identifying and isolating it from the hand.

#### 3.2.3. Histogram of Oriented Gradients (HOG)

The Histogram of Oriented Gradients (HOG) is a feature descriptor that captures the distribution of gradient orientations within localized regions of the image. By dividing the image into small cells and computing histograms of gradient directions, HOG provides a robust representation of the palm's shape and structure [2]. This technique is beneficial for palm segmentation as it emphasizes the edges and contours of the palm, making them more distinguishable from




Figure 2. Example of preprocessing steps performed on an input palm image.

```
{
  "image": "FEMALE/IMG_0025.JPG",
  "scores": {
    "luck": 0.9,
    "health": 0.7,
    "romantic": 0.4.
  }
}
```

Figure 3. An example of labels produced by a VLM, in JSON format.




Figure 4. Our proposed pipeline structure.

the background and the rest of the hand. HOG features are also less sensitive to variations in lighting and small deformations, which enhances the segmentation accuracy and reliability in real-world scenarios.

### 3.3. Data Annotation

In our study, we utilize a Vision-Language Model (VLM) to perform the initial labeling of palm line images. The VLM leverages its ability to understand and process visual and textual information simultaneously, enabling it to generate accurate labels for the intricate patterns found in palm lines. This automated labeling process serves as a foundational step in our data annotation pipeline, significantly reducing the manual effort required and ensuring consistency across the dataset. By employing a VLM, we capitalize on advanced machine learning techniques to efficiently and effectively annotate our dataset, facilitating subsequent analysis and model training. This approach not only enhances the quality of the initial labels, but also accelerates the overall data preparation process.




Figure 5. Feature Extraction with Hybrid Transformer.

## 4. Technical Approach and Model Architecture

Our system is designed as a novel, multi-modal pipeline that integrates computer vision and natural language processing to translate palm features into insightful descriptions. The architecture moves beyond a simple feature-to-text model by incorporating a rule-based scoring mechanism and a final Large Language Model (LLM) synthesis step for coherent and contextualized output.

### 4.1. Feature Extraction

#### 4.1.1. Preprocessing

The pipeline begins by preprocessing the raw palm image to standardize input and highlight regions of interest. This involves:

- **Color Correction & Normalization:** Adjusting for lighting variations to ensure consistent color and contrast across images.
- **Pale Region Isolation:** A specialized segmentation model (e.g. a lightweight U-Net) is employed to accurately identify and crop the four primary pale regions corresponding to the mounts of the palm (e.g., Mount of Venus, Mount of Jupiter, etc.), as well as the central palmar area containing the primary lines (Heart, Head, Life).
- **Image Patching:** The isolated full palm and/or specific regions are divided into a sequence of fixed-size patches, which are then linearly projected into a lower-dimensional space to create a patch embedding sequence.

#### 4.1.2. Feature Extraction with Hybrid Transformer

Instead of a single transformer, we propose a Hybrid Transformer Encoder architecture with the following features, as shown in Figure 6.

- **Shared Backbone:** A Vision Transformer (ViT) backbone first processes the entire preprocessed palm image to generate a global context embedding, represented by the [CLS] token.

- **Region-Specific Encoders:** In parallel, the cropped image regions containing specific palm lines (Heart, Head, Life, Fate) are processed by smaller, independent transformer encoders. This allows for focused feature extraction tailored to the unique characteristics of each line.
- **Feature Fusion:** The global [CLS] token embedding and the features from the region-specific encoders are concatenated and passed through a fusion module (e.g., a cross-attention layer or a simple fully connected layer) to create a unified, comprehensive feature representation  $F_{\text{unified}}$ .

## 4.2. Feature Processing and Multi-Modal Interpretation

This stage interprets the fused features  $F_{\text{unified}}$  through two parallel paths.

### 4.2.1. Text Decoder

A pre-trained text decoder (e.g., a Transformer Decoder or a GPT-2 architecture) takes  $F_{\text{unified}}$  as input. It is trained to generate structured, factual descriptions in natural language for each palm line in a template-like format. Our template is defined as the following.

- **Input:**  $F_{\text{unified}}$  (Feature vector for a specific palm).
- **Expected Output Format:** "The heart line is long and curvy. The head line is deep and long. The life line curves completely around the thumb."

### 4.2.2. Embeddings and Score Prediction

This branch translates the image features into a quantitative and comparable format.

- **Embedding Generation:** The feature vector  $F_{\text{unified}}$  is projected into a lower-dimensional, normalized embedding space (e.g., using a dedicated Embedding Layer). The core hypothesis is that palms with similar palmistry traits will lie close to each other in this space (high cosine similarity). We train our model using contrastive methods, by minimizing the dot product between embeddings of palms with similar traits, and maximizing that for embeddings of palms with traits that differ more.
- **Rule-Based Scoring Engine:** This is a novel, explicit component. A set of deterministic, human-defined rules based on traditional palmistry (e.g., "if the life line is curved, add 5 points to vitality score") is applied directly to the structured descriptions generated by the Text Decoder branch. This generates a score vector  $S_{\text{rule}}$  (e.g. [Vitality: 8, Creativity: 5, Logic: 3]).
- **Neural Network Score Predictor:** In parallel, a shallow Neural Network (NN) processes the unified embedding to predict a complementary score vector  $S_{\text{NN}}$ . The model is trained so that  $S_{\text{NN}}$  correlates with the rule-based scores  $S_{\text{rule}}$  while also learning subtle, non-linear patterns from the data.

## 4.3. LLM Inference for Holistic Synthesis

The final stage leverages a Large Language Model (LLM) as a sophisticated information synthesizer.

- **Input Prompt Construction:** The outputs from previous stages are assembled into a detailed prompt for the LLM. We include structured descriptions from the text decoder, score vectors  $S_{\text{NN}}$  and  $S_{\text{rule}}$ , as well as a system prompt framing the task.
- **Natural Language Generation:** The LLM (e.g., a fine-tuned Llama 2 or GPT-3.5 model) processes this rich, multi-modal input to produce the final output. It does not just paraphrase; it reasons about the interplay between different lines and scores, generating a novel, insightful, and human-readable report that includes character analysis and life suggestions. This step adds a layer of interpretation and narrative fluency that is absent in simpler sequence-to-sequence models.

## 5. Evaluation Target

We evaluate our model on the following metrics targeting different components of our pipeline.

- **Palm Line Detection:** We set the baseline for palm line detection at  $\geq 0.85$  for mean Intersection over Union (mIoU); and an F1 score of  $\geq 0.9$  on a test set.
- **Scoring System:** We align the scores  $S_{\text{rule}}$  and  $S_{\text{NN}}$  from our rule-based scoring and neural network, and aim for a mean squared error (MSE) of  $\leq 0.15$ .
- **LLM Inference Quality:** We employ the LLM-as-judge technique, and use a large language model to label text inference as 'Pass' or 'Fail'. We aim to have at least 90% of outputs achieving 'Pass'.

## 6. Results

### 6.1. User Interface

We provide a graphical user interface for seamless use of our pipeline, through a web app. The user first scans their palm using an on-device camera, and will receive predictions in natural language with an LLM. The user may ask further questions and through a chat-like dialogue.

### 6.2. Training Results

?????

## 7. Conclusion

Palmreader.ai establishes a multi-modal pipeline that performs fortune-telling through palm reading, incorporating Vision Transformers (ViT), a Neural Network for scoring, as well as a Large Language Model (LLM) for final inference. We have provided a first step in using a data-driven approach to draw superstitious practices towards science.

You are a palmistry expert.  
You will assist the user in inferring fortune-telling results based on their palm line features.

## Supplementary Information on Palmistry  
Palm lines of interest include the \*Life Line\*, the \*Heart Line\*, the \*Fate Line\*, and the \*Head Line\*. These lines represent respectively, \*enthusiasm and strength\*, \*romantic life\*, \*fortune and luck\*, and \*smartness and potential\*. You will receive scores predicted based on palm line features by a Deep Learning model. You are then to provide fortune-telling results in natural language, further instructions will be given in the next section.

### About Each Palm Line  
<supporting information>

## Inference Instructions

You will be provided with scores for 'strength', 'romantic', 'luck', and 'potential'. These scores were previously predicted using a deep learning model based on geometric features of the user's palm lines, and they correspond to the above points provided to you as supplementary information on palmistry.

Provide the user with fortune-telling results based on these scores, and give explanations to assist the user in inferring their results. Explain what those scores mean and how their luck or fortune looks like. Give advice to the user on how to face probable events in their future life.

Figure 6. Contextual instructions provided through the system prompt to the LLM.

We expect to produce a robust model that can effectively generate accurate palm-reading results.

Further work may optimize the pipeline by employing more powerful pretrained text decoders. Additionally, external data related to palmistry may be used for Retrieval-Augmented Generation (RAG) with the LLM in the final inference step, to provide context whenever required as additional support.




Figure 7. Homepage of our web app.




Figure 8. User interface prompting for palm scanning.




Figure 9. Follow-up dialogue with LLM.

## References

- [1] Feyisayo Amujoyegbe. Human palm images. 2022.
- [2] Daniel Louise M Parulan, Jon Neil P Borcelis, and Noel B Linsangan. Palm lines recognition using dynamic image segmentation. In *2024 IEEE International Conference on Automatic Control and Intelligent Systems (I2CACIS)*, pages 238–243. IEEE, 2024.
- [3] Big Brothers Big Sisters. Palm reading. *unknown*, 2018.